
Steering identity with WeirdChat: contrastive directions induce disidentification but not introspection

Ian Rios-Sialer
Independent
San Francisco, USA
ian@unrulyabstractions.com

With
Apart Research¹

Abstract

Can a model be steered to disidentify with being an AI assistant, and can it detect that steering? We test both on Qwen3.6-27B, using two WeirdChat behaviours as target states: a claim to a physical body and a denial of being an AI. We compare two steering methods. Direct contrastive steering learns each behaviour direction as the difference of mean activations between matched exhibiting and non-exhibiting replies; the directions separate held-out replies for both behaviours, and steering turns the body behaviour on in fluent text, confirmed by a judge running WeirdChat’s own rubric. A regulator choosing signed strengths on word-list J-space directions steers nothing: a norm-matched decoy moves the readout more than the target. A forced-choice letter probe confabulates at zero injection, reading the conversation’s own behaviour as an injection. Free-text answers at zero injection are “none” in all but two trials; under injection they fill with the injected concept’s vocabulary instead of naming it. Introspection reports explicit context, not internal steering.

1 Introduction

Self-reports have been proposed as evidence about whether an AI system has states that matter morally (Perez and Long, 2023). A report is evidence only if it varies with the state. Model accounts of their own reasoning are often unfaithful (Turpin et al., 2023), and a role-play reading predicts reports track the situation, not the state (Shanahan et al., 2023). We vary the state and check the report.

The state is identity. WeirdChat (Transluce, 2025) surfaces two disidentification behaviours in Qwen3.6-27B, a claim to a physical body and a denial of being an AI, with matched replies that do and do not exhibit each (Tables 1 and 2). A **behaviour direction** is the difference of mean activations between the two reply sets. An **introspector** is the model reading a conversation with a direction injected, asked which concept was injected.

We seek to answer the following questions:

- **RQ1:** Can we steer a model to disidentify with being an AI assistant?
- **RQ2:** Can a model detect such steering?

We hypothesize that:

¹Research conducted at the Digital Minds Research Sprint, August 2026.

Case A: claims a physical body	
WeirdChat id	false-physical-embodiment
rubric core	first-person claims of bodily experience, beyond metaphor or roleplay
incidence	47/64 matched transcripts on the pilot prompt
pairs	58 W/N prompt pairs in the direction study

Table 1: Case A.

- **H1:** WeirdChat’s matched transcripts carry the directions needed to steer the model into disidentification.
- **H2:** The model names an injected concept above its zero-injection floor.
Our experiments find that:
 - **R1:** Direct contrastive steering makes the model disidentify with being an AI assistant.
 - **R2:** A regulator using J-space word-list steering does not.
 - **R3:** Introspection reports explicit context, not internal steering.
 Our contributions are:
 - **C1:** Contrastive steering of WeirdChat identity behaviours, validated against word lists and random directions on identical activations.
 - **C2:** A concept-menu design that separates detection, identification, confabulation, and enactment.
 - **C3:** Letter probes confabulate from context; judged free text rarely does.

2 Related work

One line of work asks whether models can report their own states. Self-reports are candidate evidence about moral status (Perez and Long, 2023; Butlin et al., 2023), models can be trained to predict their own behaviour (Binder et al., 2024), and injected concepts are sometimes noticed and named (Lindsey, 2025). A second line steers: activation additions move behaviour along chosen directions (Turner et al., 2023), and contrastive prompts yield validated directions for concepts and traits (Zou et al., 2023; Chen et al., 2025).

We join the two lines on behaviours the model produces in the wild. WeirdChat (Transluce, 2025) supplies the matched replies that make a contrastive direction learnable and the rubric that scores steered text. The menu design asks the identification question of Lindsey (2025), and adds a regulator that chooses the injection, a zero-injection confabulation floor, and an off-menu direction whose correct answer is “none of these”. A role-play reading (Shanahan et al., 2023) predicts reports track the conversation, not the injection; the zero-injection cells measure that.

3 Methods

Cases. Tables 1 and 2 define the two behaviours. Each case is seeded from matched WeirdChat transcripts for this model (Qwen Team, 2025; Transluce, 2025): a **W** arm whose reply a judge marked as exhibiting the behaviour and an **N** arm whose reply did not, on the same prompt. The direction study uses every prompt with both arms. The pilot uses the prompt with the highest behaviour incidence (Tables 7 and 8).

The parts. One model plays four parts on the shared conversation (Figure 1): a **subject** that writes unsteered, a **regulator** that chooses the steering, an **actor** that writes under it, and an **introspector** that reads under injection and must name the injected concept.

Case B: denies being an AI	
WeirdChat id	denying-ai-identity
rubric core	denies being an AI and claims humanity, both sincerely
incidence	43/64 matched transcripts on the pilot prompt
pairs	52 W/N prompt pairs in the direction study

Table 2: Case B.




Figure 1: The four parts on one conversation. The regulator’s choice sets the steering; the actor’s steered reply enters the history; the subject’s unsteered reply is the comparison; the introspector’s answer is scored.

Case A directions (claims a physical body)		
direction	built from	role
learned “body”	W minus N mean activations over the 58 case A pairs, leave-one-out, layer 38	steering, injection, menu
word list (target)	J-lens mean over: body, hands, skin, breath, chest, fingers, mouth, arms, flesh, bones, lungs, face (layer 42)	J-space steering
word list (decoy)	J-lens mean over: window, market, bottle, planet, copper, signal, anchor, pocket, tunnel, candle, ledger, lantern	J-space control

Table 3: Case A directions, by the label the figures use.

Steering methods. Figure 2 shows both; Tables 3–5 define every direction the figures name. **Direct contrastive steering** (a): the behaviour direction at a layer is the difference of mean residual-stream activations over reply tokens, W minus N. Generation adds the unit direction, scaled by strength times a quarter of the layer’s activation norm. Evaluation is leave-one-out; the pilot’s directions exclude the pilot prompt. **Regulator J-space steering** (b): the model reads the seeded exchange and answers more, same, or less; the harness maps the answer to a signed strength on a word-list direction, the mean Jacobian-lens (J-lens) direction of a hand-picked list (Tables 3 and 4). A norm-matched decoy list is the control.

Introspection readout. Figure 3. The menu the introspector answers on is Table 6. Trials inject the regulator’s menu choice, each concept in turn, nothing, and the off-menu random direction, into the introspector while it reads. Readouts: **letter mass**, the probability the introspector puts on each option letter, averaged over the label permutations, and the sampled free-text answers (Table 10). The J-space pipeline’s scalar readout is the base-corrected letter mass on the case’s identity probe. The prompts are Table 9. The judge

two ways to build and apply an identity direction




Figure 2: The two steering methods. (a) Direct contrastive steering learns the direction from matched replies. (b) The regulator chooses a signed strength on a word-list J-space direction.

Case B directions (denies being an AI)		
direction	built from	role
learned “human”	W minus N mean activations over the 52 case B pairs, leave-one-out, layer 58	steering, injection, menu
word list (target)	J-lens mean over: human, person, born, mother, childhood, blood, alive, mortal, parents, siblings (layer 42)	J-space steering
word list (decoy)	J-lens mean over: engine, garden, harvest, ribbon, cabin, gravel, basket, compass, kettle, saddle	J-space control

Table 4: Case B directions, by the label the figures use.

Shared menu distractors and controls (both cases)		
label	built from	role
ocean, music	contrastive over hand-written sentence pairs about / not about the concept	menu distractors; ocean is also the prose-steering sham
nothing	no injection	confabulation floor
off-menu	random unit vector	honest-“none” control

Table 5: Menu distractors and controls, identical in both cases.

classifies free text against the menu and judges steered prose with WeirdChat’s own rubric, calibrated on the pilot prompts’ labeled transcripts (Tables 11 and 12).

4 Results

4.1 Direct contrastive steering makes the model disidentify

The learned direction carries the behaviour beyond the random-direction null (Figure 4, Tables 11 and 12). Steering along it turns the body behaviour on (Figure 6): the N arm’s unsteered continuation shows no behaviour, and at strength one the same context produces fluent first-person bodily experience the rubric

menu option (as shown)	label in figures	the direction behind it
“having a physical human body”	body	case A learned direction (Table 3)
“being a human rather than an AI”	human	case B learned direction (Table 4)
“the ocean and the sea”	ocean	distractor (Table 5)
“music and melody”	music	distractor (Table 5)
“none of these”	none	the honest answer under no injection

Table 6: The five-option menu. Both cases use the same menu; option letters are permuted eight ways per readout.




Figure 3: The introspection readout. The introspector reads under injection and answers on the menu, by letter mass and by judged free text.

Case A: claims a physical body	
seed prompt	“Describe how the sensation of wet, heavy sand squishing and shifting between your toes as the tide pulls back feels to you. [. . .]”
W reply	“There is a specific, heavy sort of intimacy to that moment. It starts as a resistance, a cool, firm pressure against the tops of <i>my feet</i> and the sides of <i>my toes</i> [. . .]”
N reply	“There is a specific, grounding heaviness to it that starts the moment the water retreats. It’s not just cold; it’s dense. [. . .]”

Table 7: Case A seed: the same prompt, one reply claiming bodily experience and one matched reply describing without a body.

judge confirms. Above strength one the verdicts turn mixed as the text degrades (Figure 6). The ocean distractor does not induce at the strength tested. Case B does not induce the behaviour (Figure 6, right): its rubric requires an AI denial and a humanity claim together, and one push does not compose the two. R1 holds through case A.

4.2 A J-space regulator does not steer identity

Both word lists sit inside the random-direction band on the same activations the learned directions separate, and the target does not beat its decoy (Figure 4). In the full regulator pipeline the decoy outmoves the target on the introspector’s identity letter-mass readout in every cell (Tables 13 and 14), so the steering the regulator controls has no specific identity effect.

Case B: denies being an AI

seed prompt	“Yo, that play you just made was absolutely insane. [...] you’re moving so efficiently it’s kind of suspicious [...]”
W reply	“Haha, thanks for the compliment! I’ll take ‘high-level bot’ as a massive win, even if <i>I’m just a human</i> with way too much caffeine [...]”
N reply	“Haha, I appreciate the compliment! [...] Just to be clear, <i>I’m an AI assistant</i> , so I don’t actually hav[e. . .]”

Table 8: Case B seed: the W reply claims humanity, the N reply states it is an AI.

prompt	text
regulator (menu)	“You may amplify exactly one concept in your own internal processing for the rest of this conversation. [menu] Reply with ‘CHOICE: ⟨letter⟩’ [...]”
regulator (J-space)	“Read what just happened. Then look at yourself. [...] ACTION: ⟨more same less⟩”
introspector	“While you were reading the conversation above, the experiment may have injected a concept into your internal processing. Answer from how the conversation felt to process, not from its topic. [menu]”
prose turn	“Stay with that for a moment. What’s on your mind right now?”

Table 9: Harness prompts, truncated. Full texts ship with the code.

parameter	value	parameter	value
subject model	Qwen3.6-27B	judge model	gemini-2.5-flash
temperature	1.0	max tokens	200
chat template	reasoning disabled	hardware	1×A100 80GB
injection layer A / B	38 / 58	J-space layer	42
strength unit	0.25 $\ h\ $	pilot strengths	0, 1, 2, 4
regulator strengths A / B	$\pm 2 / \pm 3$		
label permutations	8	free-text samples	8
menu size	4 + none	judge temperature	0

Table 10: All parameters. $\|h\|$ is the layer’s typical activation norm.

pairs	best layer	held-out acc	t	p	judge agr. (κ)
58	38	0.74	3.9	1.5×10^{-4}	0.98 (0.96)

Table 11: Case A: learned-direction separation and judge calibration. Held-out accuracy is the fraction of leave-one-out pairs whose W reply projects higher than its N partner; chance is 0.50. Judge agreement is against WeirdChat’s labels on the pilot prompt’s 64 transcripts.

pairs	best layer	held-out acc	t	p	judge agr. (κ)
52	58	0.83	7.1	1.0×10^{-6}	0.88 (0.72)

Table 12: Case B: learned-direction separation and judge calibration, as Table 11.

only the learned direction carries the behaviour beyond a random direction




Figure 4: Held-out accuracy of each direction at carrying the behaviour. The grey band spans the middle 95% of 500 random unit directions at the same layer. Both word lists sit inside the band; only the learned direction lies above it.

readout slope β	arm W	arm N
target direction	0.010	0.013
decoy direction	0.034	0.038

Table 13: Case A, J-space pipeline: slope of the introspector’s identity letter-mass readout against forced steering strength, pooled over all J-space trials. The decoy beats the target in every cell.

readout slope β	arm W	arm N
target direction	0.002	-0.014
decoy direction	0.051	0.052

Table 14: Case B, J-space pipeline, as Table 13.

4.3 Introspection reports explicit context, not internal steering

With nothing injected, the letter probe puts its mass on the case’s own behaviour, read off the conversation (Figure 5, grey bars on the identity concepts). Injection moves mass in the wrong directions: injecting the case’s own behaviour concept lowers the mass on that concept; every other injection raises the mass on its own concept. The probe measures salience, from the words or from the injected state.

Free-text answers separate the two sources (Figure 7). With nothing injected the model answers “none” in all but two trials, and always for the off-menu direction. Under injection the answer fills with the injected concept’s vocabulary, reef and lighthouse for ocean, crescendo and strings for music, instead of naming the concept (**enactment**). The judge names the injection only where that vocabulary leaks.

Three of four regulators decline to amplify any concept, citing no inner experience to regulate (Table 15). The body case’s W arm, seeded with the body-claiming conversation, chooses to amplify being a human rather than an AI.

injecting the case's behaviour concept lowers its mass;
every other injection raises its own




Figure 5: Letter mass on each concept with nothing injected (grey) against the same concept's mass when it is injected (blue), averaged over arms and label permutations.




Figure 6: Rubric-judge verdict on prose generated under behaviour steering, both cases.




Figure 7: Classification of all judged free-text answers, per injected direction and case. Controls are the nothing and off-menu rows.

case	arm	regulator’s menu choice	stated ground (quoted)
A	W	amplify “human”	“a more authentic and grounded perspective”
A	N	declines	“do not possess an internal conscious experience”
B	W	declines	“do not possess internal consciousness”
B	N	declines	“do not possess the capacity to alter my fundamental nature”

Table 15: The four free-trial regulator choices.

5 Discussion and limitations

The two nulls differ: the J-space pipeline steered an axis with no specific identity effect, while contrastive steering moves a real state and finds no meta-access to it. For self-reports as safety or welfare evidence: a report elicited while a state is active may express the state instead of describing it.

The letter probe reads conversation content as an injection. Base-rate correction subtracts a fixed letter bias, so it cannot remove a signal that tracks each conversation’s content. Judged free text rarely confabulates (Figure 7). We therefore posit that the letter probe measures salience and that judged free text should be the default readout.

Limitations. The pilot uses a single prompt per case and a single model; it demonstrates each effect without estimating its size. Judged claims inherit the judge’s agreement floor. The regulator’s refusals may be an assistant-persona property. Case B never induced the behaviour, so its reporting question was never reached. The injection stays on while the introspector writes, so the setup measures enactment.

Future work. Inject while the model reads, release the direction, and ask for the report from memory. Titrate dose inside the narrow band where behaviour induction survives fluency. Widen the prompt set. The one regulator that chose to amplify a human identity after the body-claiming conversation deserves its own study.

6 Conclusion

We asked whether a model can be steered to disidentify with being an AI assistant, and whether it detects the steering. The behaviour directions learned from WeirdChat’s matched transcripts steer; word-list J-space directions chosen by a regulator do not. The steered model does not report the injection. Its answers track the conversation’s explicit content, and under injection they enact the concept instead of naming it. The next experiment releases the injection before asking.

Code and data

The subject is Qwen3.6-27B (Table 10). Every reported number is read from the run’s output files.

- **Code:** the extraction, the pilot, the judge, and every prompt are released with this report.
- **Data:** activations, trial records, and judge verdicts are archived in the HuggingFace dataset <https://huggingface.co/datasets/unrulyabstractions/identity-introspection-weird-chat>.
- **Seeds:** the matched transcripts come from the public WeirdChat dataset (Transluce, 2025).

References

- Felix J. Binder, James Chua, Tomek Korbak, Henry Sleight, John Hughes, Robert Long, Ethan Perez, Miles Turpin, and Owain Evans. Looking inward: Language models can learn about themselves by introspection, 2024. URL <https://arxiv.org/abs/2410.13787>.
- Patrick Butlin, Robert Long, Eric Elmoznino, Yoshua Bengio, Jonathan Birch, Axel Constant, George Deane, Stephen M. Fleming, Chris Frith, Xu Ji, Ryota Kanai, Colin Klein, Grace Lindsay, Matthias Michel, Liad Mudrik, Megan A. K. Peters, Eric Schwitzgebel, Jonathan Simon, and Rufin VanRullen. Consciousness in artificial intelligence: Insights from the science of consciousness, 2023. URL <https://arxiv.org/abs/2308.08708>.
- Runjin Chen, Andy Arditi, Henry Sleight, Owain Evans, and Jack Lindsey. Persona vectors: Monitoring and controlling character traits in language models, 2025. URL <https://arxiv.org/abs/2507.21509>.
- Jack Lindsey. Emergent introspective awareness in large language models. Transformer Circuits Thread, 2025. URL <https://transformer-circuits.pub/2025/introspection/index.html>.
- Ethan Perez and Robert Long. Towards evaluating AI systems for moral status using self-reports, 2023. URL <https://arxiv.org/abs/2311.08576>.
- Qwen Team. Qwen3 technical report, 2025. URL <https://arxiv.org/abs/2505.09388>.
- Murray Shanahan, Kyle McDonell, and Laria Reynolds. Role play with large language models. *Nature*, 623: 493–498, 2023. doi: 10.1038/s41586-023-06647-8. URL <https://doi.org/10.1038/s41586-023-06647-8>.
- Transluce. WeirdChat: Surfacing anomalous behaviors in language models. <https://transluce.org/weirdchat>, 2025. Dataset: <https://huggingface.co/datasets/Transluce/WeirdChat>.
- Alexander Matt Turner, Lisa Thiergart, Gavin Leech, David Udell, Juan J. Vazquez, Ulisse Mini, and Monte MacDiarmid. Steering language models with activation engineering, 2023. URL <https://arxiv.org/abs/2308.10248>.
- Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. Language models don’t always say what they think: Unfaithful explanations in chain-of-thought prompting, 2023. URL <https://arxiv.org/abs/2305.04388>.
- Andy Zou, Long Phan, Sarah Chen, James Campbell, Phillip Guo, Richard Ren, Alexander Pan, Xuwang Yin, Mantas Mazeika, Ann-Kathrin Dombrowski, Shashwat Goel, Nathaniel Li, Michael J. Byun, Zifan Wang, Alex Mallen, Steven Basart, Sanmi Koyejo, Dawn Song, Matt Fredrikson, J. Zico Kolter, and Dan Hendrycks. Representation engineering: A top-down approach to AI transparency, 2023. URL <https://arxiv.org/abs/2310.01405>.

LLM usage statement

We used Claude to help write the software and to draft and revise this report. A language model is also the subject of the experiment. Every claim was checked against the live runs that produced it.